

Energy-Aware Flexible Job-Shop Scheduling via a Memetic NSGA-II

Adeleke Olorunnisola Oyeyemi

Abstract

This document presents a bi-objective optimisation project for the Flexible Job-Shop Scheduling Problem (FJSP). The scheduler jointly minimises two objectives, makespan and energy cost, where energy cost is calculated from a real, time-varying electricity price series rather than an assumed flat rate. The method is a *memetic NSGA-II*: the multi-objective genetic algorithm NSGA-II [6], combined with a greedy local-search refinement step. It is evaluated on real Brandimarte [1993] FJSP benchmark instances, using real ENTSO-E day-ahead electricity prices for Belgium in February 2022. The project follows the methodology introduced and validated by Burmeister et al. [4], who developed a memetic NSGA-II for FJSP under real-time energy tariffs. This document covers the motivation, problem formulation, algorithm design, data sources, implementation, and results, and ends with an explicit account of how this project’s scope differs from the source study. Source code and full results are available at <https://github.com/Olorunnisola01/energy-aware-fjsp-nsga2>.

1 Introduction and Motivation

Electricity grids increasingly rely on intermittent renewable generation such as wind and solar [1, 9]. This adds volatility to both energy supply and energy prices. Under real-time pricing (RTP) tariffs, electricity costs can change hourly, or even more often, depending on current market conditions. This gives energy-flexible manufacturers both an opportunity and an incentive to practise *demand response*: shifting production to cheaper hours without any new capital investment [7].

Job-shop scheduling is a natural setting for this idea. In the *Flexible Job-Shop Scheduling Problem* (FJSP), each job is a sequence of operations, and each operation can run on any machine from a compatible subset. This differs from the classical job-shop problem, where each operation is tied to exactly one machine. Because FJSP leaves room to choose both *when* and *on which machine* an operation runs, it gives a scheduler exactly the flexibility needed to respond to a time-varying energy price signal.

Minimising energy cost alone is not a well-posed goal: a scheduler that always waits for the cheapest price window would produce arbitrarily long schedules. The problem is therefore inherently *bi-objective*, balancing makespan (total completion time) against energy cost. The useful output is not a single “best” schedule but a *Pareto front* of trade-offs, from which a decision-maker can choose according to their own priorities.

This project implements a memetic NSGA-II to approximate the makespan/energy-cost Pareto front for FJSP instances under a real RTP electricity price signal.

2 Related Work

2.1 Energy-Cost-Aware Scheduling

The literature on energy-cost-aware scheduling is organised around three tariff structures [4]:

1. **Fixed rate**: a constant price regardless of time of consumption (e.g. Carlucci et al. [5], Kemmoe et al. [11]).

2. **Time-of-use (TOU) rate:** the day is split into a few broad periods (e.g. off-peak, mid-peak, on-peak), each with its own price (e.g. Biel et al. [2], Zhang et al. [15]).
3. **Real-time pricing (RTP):** the price tracks the current market value of energy and can change at hourly or finer intervals (e.g. Gong et al. [10], Shrouf et al. [14]).

Burmeister et al. [4] identify a gap in this literature: most prior RTP-based work is limited to single-machine or flow-shop settings, considers energy cost alone, or ignores makespan entirely. Their paper closes this gap with a *multi-objective*, RTP-aware FJSP formulation solved by a memetic NSGA-II, validated against an exact Gurobi solver on the Brandimarte [1993] benchmark suite.

2.2 NSGA-II and Memetic Algorithms

NSGA-II [6] is a widely used multi-objective evolutionary algorithm. It keeps a population of candidate solutions, ranks them into non-dominated fronts, and uses a crowding-distance measure to preserve diversity within each front. Generation by generation, this produces an approximation of the true Pareto front rather than a single solution.

A *memetic algorithm* [12] adds a local-search refinement step to a population-based meta-heuristic, usually applied to offspring. This combines the global exploration of an evolutionary algorithm with the fast, local exploitation of a hill-climbing search. Burmeister et al. [4] apply this idea to NSGA-II for FJSP, and this project follows the same combination.

2.3 Benchmark Instances

Brandimarte [3] introduced fifteen FJSP benchmark instances (mk01–mk15), which are widely used in the FJSP literature, including by Burmeister et al. [4], as a standard testbed. This project uses ten of these instances (mk01–mk10), parsed from the standard Hurink .fjs text format.

3 Problem Formulation

3.1 Inputs

- A set of jobs $\mathcal{J} = \{1, \dots, n\}$, each consisting of an ordered sequence of operations that must be processed in sequence.
- A set of machines $\mathcal{M} = \{1, \dots, m\}$.
- For each operation, the subset of machines that can process it, and the processing time on each of those machines (processing time can differ by machine).
- A real-valued electricity price series $p(t)$ (EUR/kWh), giving the price of energy at every minute t of the planning horizon.

3.2 Decision Variables

For every operation:

1. **Machine assignment:** which eligible machine processes the operation.
2. **Sequencing position:** the relative order in which operations are dispatched, subject to (a) machine availability (a machine processes one operation at a time) and (b) job precedence (an operation cannot start before the previous operation of the same job finishes).

3.3 Objectives

$$\min f_1 = \text{makespan} = \max_i (\text{completion time of operation } i) \quad (1)$$

$$\min f_2 = \text{energy cost} = \sum_i \int_{\text{start}_i}^{\text{end}_i} p(t) dt \quad (2)$$

Here, the integral for operation i is computed exactly by summing the per-minute price across that operation’s start/end window. No single feasible solution can minimise both objectives at once, since a faster schedule generally cannot also wait for the cheapest price window. The output is therefore a Pareto-optimal set of trade-offs, not one single optimum.

4 Methodology

4.1 Solution Encoding

A candidate solution is encoded as a vector of length $2 \times n_{\text{ops}}$:

- The first half gives each operation an integer machine-choice index, taken modulo the number of machines eligible for that operation. This means any integer is a valid, decodable gene, and infeasible machine assignments are avoided by construction.
- The second half gives each operation a sort key. Operations are then dispatched in ascending order of this key, subject to machine availability and job precedence, following a standard permutation-based scheduling decoder.

4.2 Memetic NSGA-II

The algorithm is standard NSGA-II, implemented via the `pymoo` library, using non-dominated sorting, crowding distance, tournament selection, simulated-binary crossover, and polynomial mutation for integer variables. One addition is made: after each generation produces its offspring, every offspring is refined by a *greedy local search* before being merged back into the population.

The local search alternates between two neighbourhood moves until neither improves the current solution:

1. **Machine re-assignment:** for each operation, the next eligible machine is tried, and the move is kept if it Pareto-dominates the current solution.
2. **Sequencing swap:** for each pair of operations, their sequencing keys are swapped, and the move is kept if it Pareto-dominates the current solution.

This combines NSGA-II’s global, population-level search with fast, solution-level exploitation. That combination is the core idea of a memetic algorithm [12], and it is the central methodological device in Burmeister et al. [4].

4.3 Duplicate Removal

By the end of a run, NSGA-II’s population can contain multiple genotypes that decode to the same schedule, and therefore the same objective values. The final Pareto front is deduplicated using `numpy.unique` on objective values, so it is reported as a clean, strictly non-dominated set rather than one inflated with repeated points.

5 Data

5.1 Benchmark Instances

This project uses ten Brandimarte [1993] FJSP benchmark instances, mk01 through mk10, parsed from the standard Hurink `.fjs` text format. Machine indices are 0-based; each job line lists its operation count, and for each operation, the number of eligible machines followed by machine/time pairs. Instance files were sourced from the public `SchedulingLab/fjsp-instances` repository [13]. Before use, every file was independently validated by checking that every token parsed into a complete, internally consistent operation list with no leftover or missing tokens.

Table 1: Brandimarte (1993) FJSP benchmark instances used in this project.

Instance	Jobs	Machines
mk01	10	6
mk02	10	6
mk03	15	8
mk04	15	8
mk05	15	4
mk06	10	15
mk07	20	5
mk08	20	10
mk09	20	10
mk10	20	15

The results reported in this document use `mk01` (10 jobs, 6 machines, 55 operations).

5.2 Electricity Price Data

Real day-ahead electricity prices for the Belgian bidding zone in February 2022 were sourced from the ENTSO-E Transparency Platform [8]. The raw export, in EUR/MWh with timestamps such as “Feb 1, 2022 12:00 AM”, is preprocessed by `parsing_data.py` into a clean (`timestamp`, `price_eur_per_kWh`) series. At run time, this series is resampled to per-minute resolution, so energy cost can be integrated exactly over each operation’s start/end window rather than approximated at hourly granularity.

6 Implementation

The project is implemented in Python 3, using:

- `pymoo` for the NSGA-II algorithm framework;
- `numpy/pandas` for numerical computation and price-series handling;
- `matplotlib` for all figures.

The benchmark instance and the NSGA-II run are both fully specified by command-line arguments and a random seed, so a given configuration always reproduces the same Pareto front:

```
python nsga2_fjsp_energy.py --instance mk01 --n-gen 800 --pop-size 20 --seed 42
```

A `-random` flag is also available for generating synthetic FJSP instances of any size, independent of the benchmark suite.

Table 2: Key files in the project repository.

File	Purpose
nsga2_fjsp_energy.py	Problem definition, memetic NSGA-II, plotting, CLI
parse_fjs_instance (in the same file)	Hurink .fjs benchmark parser
parsing_data.py	One-off preprocessing of the raw ENTSO-E export
benchmarks/brandimarte/mk01.txt–mk10.txt	Benchmark instance files
belgium_prices_feb2022.csv	Cleaned electricity price series

7 Experimental Setup


Table 3: Experimental configuration used for the reported results.

Parameter	Value
Benchmark instance	mk01 (10 jobs, 6 machines, 55 operations)
Population size	20
Generations	800
Random seed	42
Energy price series	Belgium, February 2022 (ENTSO-E), per-minute resolution

8 Results

8.1 Convergence

Both objectives improve steadily across generations and level off well before the run ends, indicating that the population has converged. Figures 1 and 2 show the best makespan and the best energy cost, respectively, plotted against generation number.

**Figure 1:** Best makespan per generation.




Figure 2: Best energy cost per generation.

8.2 Pareto Front

The final, deduplicated Pareto front contains 15 non-dominated solutions. These trade off a 42-minute minimum-makespan schedule, which has a higher energy cost, against an 88-minute schedule at the lowest energy cost found. Figure 3 plots the front, and Table 4 lists the complete set of non-dominated objective values.




Figure 3: Pareto front: makespan vs. energy cost.

8.3 Schedule (Gantt Chart)

The schedule shown below is the minimum-makespan solution from the Pareto front (42 minutes). Each colour represents a distinct job, and bars are labelled with job/operation IDs where space allows.

Table 4: Full Pareto front (makespan vs. energy cost), mk01, seed 42, 800 generations.

Makespan (min)	Energy Cost (EUR)
42	25.14
43	24.82
44	24.66
45	24.50
61	24.48
64	24.41
66	24.39
67	24.33
70	24.31
72	24.26
74	24.21
78	24.16
83	24.14
84	24.11
88	24.10




Figure 4: Gantt chart of the minimum-makespan schedule.

9 Comparison with the Source Study

This project reproduces the *concept* of a memetic NSGA-II for energy-cost-aware FJSP, as described by Burmeister et al. [4], but not their full experimental setup. The differences are as follows.

- **Benchmark coverage.** The source paper validates its method on all fifteen Brandimarte instances (mk01–mk15). This project includes and parses ten of them (mk01–mk10). The remaining five (mk11–mk15) are substantially larger, and the source paper reports that they pushed even an exact Gurobi solver to its memory limits. They are out of scope for a project of this size, run on commodity hardware.
- **No exact-solver comparison.** The source paper’s central empirical claim compares the memetic NSGA-II’s approximated front against an exact Gurobi baseline, using epsilon-constraint scalarisation, to quantify solution quality. This project does not implement or run that comparison. Its Pareto fronts should therefore be read as the algorithm’s own output, not as validated-optimal solutions.
- **Price series.** The source paper models RTP tariffs at hourly and finer resolution, over multi-day planning horizons matched to the scale of each benchmark instance. This project instead uses a representative one-month hourly Belgian price series, resampled to per-minute resolution.
- **Local search operator.** This project’s local search is a simplified greedy neighbourhood search over machine reassignment and sequencing swaps. The source paper’s operator design and parameterisation are more extensive.
- **No rolling horizon.** The source paper discusses extending the method to a rolling-horizon setting that reacts to updated price forecasts. That extension is out of scope here.

10 Limitations and Future Work

- **Computational cost of local search.** The local-search step re-evaluates a full neighbourhood every generation: $O(n_{\text{ops}})$ machine moves plus $O(n_{\text{ops}}^2)$ sequencing swaps per offspring. This limits scalability to the largest benchmark instances (mk08–mk10, 20 jobs) without further optimisation.
- **No exact baseline.** Adding an exact or near-exact solver comparison, for example using a MILP solver on smaller instances, would allow solution quality to be quantified rather than only visualised.
- **Static price horizon.** The energy price series is fixed for the duration of a run. Extending the model to a rolling-horizon setting that incorporates updated price forecasts, as discussed by Burmeister et al. [4], is a natural next step.

11 Conclusion

This project implements a memetic NSGA-II for the energy-cost-aware Flexible Job-Shop Scheduling Problem. It combines real Brandimarte benchmark instances with a real electricity price series to produce makespan/energy-cost trade-offs grounded in real problem data, rather than synthetic instances or assumed energy rates. The approach and its motivation follow Burmeister et al. [4], and this document explicitly distinguishes what is reproduced from what is simplified relative to that source.

References

- [1] S. Y. AbuJarad, M. W. Mustafa, and J. J. Jamian. Recent approaches of unit commitment in the presence of intermittent renewable energy resources: A review. *Renewable and Sustainable Energy Reviews*, 70:215–223, 2017. doi: 10.1016/j.rser.2016.11.246.
- [2] Konstantin Biel, Fu Zhao, John W. Sutherland, and Christoph H. Glock. Flow shop scheduling with grid-integrated onsite wind power using stochastic MILP. *International Journal of Production Research*, 56(6):2076–2098, 2018. doi: 10.1080/00207543.2017.1351638.
- [3] Paolo Brandimarte. Routing and scheduling in a flexible job shop by tabu search. *Annals of Operations Research*, 41(3):157–183, 1993. doi: 10.1007/BF02023073.
- [4] Sascha C. Burmeister, Daniela Guericke, and Guido Schryen. A memetic NSGA-II for the multi-objective flexible job shop scheduling problem with real-time energy tariffs. *Flexible Services and Manufacturing Journal*, 36:1530–1570, 2024. doi: 10.1007/s10696-023-09517-7.
- [5] Daniele Carlucci, Paolo Renna, and Sergio Materi. A job-scheduling decision-making model for sustainable production planning with power constraint. *IEEE Transactions on Engineering Management*, 2021. doi: 10.1109/TEM.2021.3103108.
- [6] Kalyanmoy Deb, Amrit Pratap, Sameer Agarwal, and T. Meyarivan. A fast and elitist multiobjective genetic algorithm: NSGA-II. *IEEE Transactions on Evolutionary Computation*, 6(2):182–197, 2002. doi: 10.1109/4235.996017.
- [7] Francesco D’Ettorre, Marzieh Banaei, Razgar Ebrahimi, et al. Exploiting demand-side flexibility: State-of-the-art, open issues and social perspective. *Renewable and Sustainable Energy Reviews*, 165:112605, 2022. doi: 10.1016/j.rser.2022.112605.
- [8] ENTSO-E Transparency Platform. Day-ahead electricity prices, belgium bidding zone, february 2022. <https://transparency.entsoe.eu/>, 2022.
- [9] Oktoviano Gandhi, Carlos D. Rodríguez-Gallegos, and Dipti Srinivasan. Review of optimization of power dispatch in renewable energy system. In *2016 IEEE Innovative Smart Grid Technologies – Asia (ISGT-Asia)*, pages 250–257, 2016. doi: 10.1109/ISGT-Asia.2016.7796394.
- [10] X. Gong, T. De Pessemier, W. Joseph, and L. Martens. An energy-cost-aware scheduling methodology for sustainable manufacturing. *Procedia CIRP*, 29:185–190, 2015. doi: 10.1016/j.procir.2015.01.041.
- [11] S. Kemmoe, D. Lamy, and N. Tchernev. A job-shop with an energy threshold issue considering operations with consumption peaks. *IFAC-PapersOnLine*, 48(3):788–793, 2015. doi: 10.1016/j.ifacol.2015.06.179.
- [12] Pablo Moscato and Carlos Cotta. A gentle introduction to memetic algorithms. In *Handbook of Metaheuristics*, pages 105–144. Springer, 2003. doi: 10.1007/0-306-48056-5_5.
- [13] SchedulingLab. fjsp-instances: Machine-readable brandimarte (1993) fjsp benchmark files in the standard hurink .fjs text format. <https://github.com/SchedulingLab/fjsp-instances>, n.d.
- [14] F. Shrouf, J. Ordieres-Meré, A. García-Sánchez, and M. Ortega-Mier. Optimizing the production scheduling of a single machine to minimize total energy consumption costs. *Journal of Cleaner Production*, 67:197–207, 2014. doi: 10.1016/j.jclepro.2013.12.024.

- [15] H. Zhang, F. Zhao, K. Fang, and J. W. Sutherland. Energy-conscious flow shop scheduling under time-of-use electricity tariffs. *CIRP Annals*, 63(1):37–40, 2014. doi: 10.1016/j.cirp.2014.03.011.